



O'REILLY®

AI Agents Bootcamp Day 2

Production-Ready Multi-Agent Systems

Nicole Koenigstein

Course Objectives: What this Course Covers

- Learn the core mental models for agent design and how they map to frameworks such as LangGraph and CrewAI
- Use LangGraph to build single-agent and multi-agent systems with typed states, tool use, memory, and conditional decision loops
- Understand and apply collaboration and communication patterns in multi-agent systems
- Learn how to transform unstructured, semi structured, and multimodal inputs into agent-ready structured data using schema-based prompting, Pydantic validation, and MCP-driven communication patterns
- Select appropriate model types and agent roles for planning, retrieval, tool use, and fast intra-agent communication
- Design memory and context strategies that support both individual agent reasoning and group-level coordination
- Apply guardrails, secure tool execution environments, and design threat-aware workflows





Course Goal: What You'll learn

Agents aren't magic. They are carefully engineered systems that apply structured workflows to complex problems. Which is why this course follows a high-level roadmap to guide you from ideation to production. This is framework independent, LangGraph is simply the concrete implementation vehicle.

- **Phase 1 (Control):** Start by building a simple, stateful agent with basic tool use.
- **Phase 2 (Intelligence):** Add intelligence by implementing structured reasoning (CoT, ReAct) Add Human-in-the-Loop to make the system smarter and safer.
- **Phase 3 (Coordination):** Scale up by designing a multi-agent system and by choosing the proper communication and collaboration patterns (Supervisor, Hierarchical, Swarm) to handle complexity. Choosing the right models for each task.
- **Phase 4 (Reliability):** For production focus on the core pillars, such as validation, guardrails and fallbacks, to ensure the system is reliable, auditable, and ready for deployment.

About Me

AI Researcher and Practitioner in Agentic Systems



Research, consulting, teaching, and direct system implementation focused on building reliable, production-ready agentic systems, with emphasis on multi-agent architectures, evaluation, safety, and long-term system behavior. Currently pursuing a PhD on self-improving AI agents and agent benchmarking.

Advised the European Commission and IOSCO on generative AI in regulated environments and serving on advisory boards for leading AI and quantitative finance conferences. Delivered 50+ invited talks and technical workshops across academia, industry, and international conferences in the past 2.5 years.

Get the Code & Course Solutions

GitHub Repo



GitHub for course code.

Solutions Link



Link to Accompanying book website for
AI Agents – The Definitive Guide
to download course solutions after the course.
(<https://ai-agents-the-definitive-guide.com/courses>)

Password: AI_Agents_Bootcamp_25_03_2026

Stay Connected

Book Website



Accompanying book website for
AI Agents – The Definitive Guide
with additional resources.

Book Discord



Discord channel for AI Agents
– The Definitive Guide

LinkedIn



Continue the Learning

Nicolepcx / ORM_AI_Agents_Bootcamp

<> Code Issues Pull requests **Discussions** Actions Projects Wiki Security Insights Settings

Category "General" has been deleted.

Manage discussion categories

Sections are a dropdown of categories. Categories have types of discussions, and discussions within them.

New section New category

Categories without section

	Announcements Updates from maintainers	 
	Ideas Share ideas for new features	 
	Q&A Answers enabled · Ask the community for help	 
	Show and tell	 

Share your ideas and projects after the course, and get your questions answered beyond the course.

Session 1

From Unstructured and Multimodal Inputs to Structured Agent-ready Data

Course Schedule

Session 1

From unstructured and multimodal inputs to structured agent-ready data (60 minutes)

- **Intro** (5 mins)
- **Setup the coding environment. (5 mins)**
- **Presentation:** Optimizing agents for structured, unstructured, and multimodal data; aligning their outputs through schema-based prompting with Pydantic validation; enabling stable multi-agent communication using the Model Context Protocol (MCP)
- **Demo:** Transforming unstructured data retrieved from an MCP server into agent-digestible structured data
- **Hands-on Exercise:** Take a free-form task description and define a minimal schema, configuration, and prompt so that an agent can consistently turn it into structured data for a downstream tool or coding workflow
- **Q&A** (5 mins)


Set up Environment

GitHub Repo



GitHub for course code:

<https://github.com/Nicolepcx/hands-on-multimodal-AI>





The Production Failure Pattern


- Agents are chained
- Outputs become inputs downstream
- Small schema drift becomes cascading failure

Why Using Validation – The Brittle JSON Trap

The Multi-Agent Reliability Tax

# of Agents (n)	Per-Agent Accuracy (p)	System Accuracy (p^n)	Error Rate
1 Agent	98%	98.0%	2.0%
3 Agents	98%	~94.1%	~5.9%
5 Agents	98%	~90.4%	~9.6%
10 Agents	98%	~81.7%	~18.3%



Error Propagation in Multi-Agent Systems



Sequential multi-agent pipeline without validation. Each handoff multiplies failure probability. System reliability decays according to the product rule.

Failure Propagation

With and without validation gates




Validation gates break multiplicative failure chains into local recoverable events.

Schema-enforced Generation

# of Agents (n)	Per-Agent Accuracy (p_eff)	System Accuracy (p_eff ⁿ)	Error Rate
1 Agent	99.8%	99.8%	0.2%
3 Agents	99.8%	~99.4%	~0.6%
5 Agents	99.8%	~99.0%	~1.0%
10 Agents	99.8%	~98.0%	~2.0%

Error Propagation in Multi-Agent Systems



Validated multi agent pipeline. Each boundary enforces structural correctness, converting multiplicative reliability decay into controlled reliability amplification.


Contracts, Tools, Reliable Execution

- Prototypes fail in production for boring reasons
- Underspecified interfaces, brittle tool calls, weak handoffs

Goal: reliability through contracts, validation, and governed execution

End to End Production

Stable execution with explicit control points




Governance and validation occur before tool execution and before propagation.

Validation Gate and Bounded Retries


- Validate output against pinned schema
- Emit explicit success or failure signals
- Increment retry count centrally
- Return structured feedback for targeted retries

Three-tier Fallback Pipeline

Three-tier fallback pipeline for agent output, progressively escalating from strict decoding to validation and emergency canonicalization.



Two Correction Loops




In a production-grade agentic system, you are essentially running two distinct types of "Correction Loops". One happens inside the model's generation window (Instructor), and one happens inside the system's workflow (LangGraph + Pydantic)

Policy Driven Contracts and Pinning

- Central policy config defines what is allowed
- Build the Pydantic schema from policy
- Pin policy per run
- Hash policy and schema for audit and reproducibility

Policy Pinning and Schema Fingerprinting




Hashing turns “which rules were active” into an auditable artifact.

MCP as Universal Tool Contract

Tools become reliable when the server enforces contracts

- Typed IO for tool inputs and outputs
- Governance at the server boundary
- Standard transport choices for local versus multi user deployments



Checkpointing and Time Travel for Fault Tolerance

- Compile graph with a checkpointer
- Resume from last stable node
- Do not replay expensive upstream reasoning after downstream failure

The Production Recipe

- Contracts first
- Validation at every boundary
- Pinned policy and schema fingerprints
- Checkpointing for recovery and cost control
- Governed tools via MCP




Key Takeaways Presentation Session 1

From unstructured inputs to structured agent-ready data

- Reasoning is probabilistic. Execution must be deterministic.
- Schemas are contracts between cognition and action.
- Validation is a reliability boundary, not formatting.
- MCP enables protocol-stable inter-agent communication.


Pulse Check

Do you see the benefit of validation?




Demo

Transforming unstructured data retrieved from an MCP server into agent-digestible structured data



Pulse Check


Has everyone set up their Notebook environment?



Hands-on

Take a free-form task description and define a minimal schema, configuration, and prompt so that an agent can consistently turn it into structured data for a downstream tool or coding workflow

- Instructions are in the Colab Notebook.



Q&A





Break

Session 2

Memory Design and Shared Context Strategies

Course Schedule

Session 2

Memory design and shared context strategies (60 minutes)

- **Presentation:** How episodic memory preserves task-level context; how procedural memory enables routines that can be reused across tasks; capturing higher-level reasoning patterns through long-term memory structures; determining what information should be stored, when memory should be updated, and how retrieval shapes subsequent decisions; checkpointing as a mechanism for preserving state
- **Demo:** A LangGraph workflow that maintains episodic and procedural memory and uses checkpoints; how memory is read and written during execution; how time travel allows stepping back to an earlier state to inspect the reasoning path
- **Hands-on Exercise:** Extend your agent by adding a simple episodic memory that stores each tool result and a procedural memory entry that records a small reusable routine
- **Q&A** (5 min)

Memory in Multi-Agent Systems

Why Memory Matters

Without memory:

- Agents are stateless
- Decisions are short sighted
- No continuity
- No accountability
- No learning

Memory is the bridge between: Single-step reasoning → Strategic multi-step systems.
Memory determines whether your system behaves like a chatbot or and organization.

The Intuition

LLMs themselves do NOT inherently remember things. You need to intentionally add memory in.





Memory as Core Component

- Short-term memory → context within a session (chain of reasoning).
- Long-term memory → persistence across sessions (historical knowledge).
- Episodic vs Semantic vs Procedural memory.




Short-Term Memory

What It Is

Context inside a single execution or session.

- Current state object
- Active plan
- Recent tool outputs
- Loop counters
- Temporary reasoning chain

Short-Term Memory: High-Level Diagram



Prevents loops and maintains coherence within a reasoning chain. Lost after execution unless persisted.

Short-Term Memory – Where to Use It

Some examples are:

- ReAct loops
- Tool execution tracking
- Multi-step reasoning chains
- Planner state transitions

Short-Term Memory – When it Fails

Short-term memory is about coherence, not persistence.

- Long-running workflows
- Cross-session continuity
- Audit or reproducibility requirements




Long-Term Memory

What It Is

Persistent memory across sessions.

- Past decisions
- Historical outputs
- Learned preferences
- Audit logs
- Past interactions

Long-Term Memory: High-Level Diagram



Cross-session continuity: Transforms stateless agents into evolving systems. Requires storage discipline and governance.

Long-Term Memory – Where to Use It

Some examples are:

- User preference tracking
- Long-running case management
- Personal AI assistants
- Multi-day research workflows
- Organizational knowledge retention

Episodic Memory – When it Fails

Long-term memory enables strategic behavior.


- High-frequency temporary state
- Rapidly changing context that should not persist
- If no validation before storage

Procedural vs Episodic vs Semantic Memory

Procedural



Agent Code



LLM

Semantic



Documentation,
Project Info

Episodic



Relevant
Code



Examples



History




Episodic Memory

What It Is

Stored experiences, you get event logs of:

- Task execution
- Tool usage
- Decisions taken
- Errors encountered

Episodic Memory: High-Level Diagram



Experience-based recall: Allows agents to analyze what happened before. Crucial for debugging, learning, and auditability.

Episodic Memory – Where to Use It

Some examples are:

- Debugging agent decisions
- Auditable AI systems
- Post-mortem analysis
- Reinforcement or self-improvement loops
- Experiment tracking

Episodic Memory – When it Fails

Episodic memory enables system introspection.

- If logs grow without summarization
- If noise dominates signal
- If no indexing or retrieval strategy




Semantic Memory

What It Is

Abstracted factual knowledge.

- Policies
- Constraints
- Domain knowledge
- Learned summaries

Semantic Memory: High-Level Diagram



Fact-level knowledge repository: Supports reasoning beyond current context window. Foundation for RAG systems.

Semantic Memory – Where to Use It

Some examples are:

- RAG pipelines
- Domain rule retrieval
- Policy enforcement
- Knowledge assistants
- Context augmentation

Semantic Memory – When it Fails

Semantic memory enables factual grounding.

- If embedding drift occurs
- If outdated information persists
- If no confidence scoring




Procedural Memory

What It Is

Encodes workflows and skills:

- Workflows
- Tool invocation logic
- Validation pipelines
- Execution policies

Procedural Memory: High-Level Diagram



Operational skill memory: Defines how tasks are executed. Changing procedural memory changes system behavior.

Procedural Memory – Where to Use It

Some examples are:

- Repeatable automation tasks
- Stable tool invocation logic
- Compliance workflows
- Deterministic pipelines
- Agentic orchestration frameworks

Procedural Memory – When it Fails

Procedural memory defines operational behavior.


- When tasks become exploratory
- When workflows are too rigid
- When adaptation is required



From Type to Topology


Per-Agent Local Memory

Each agent maintains its own history.




Isolated state is good for modular systems. Risk: agents drift with inconsistent world models.

Shared Global Memory



Single source of truth: Transparent but synchronization-heavy.

Hybrid Memory Architecture



Local autonomy + shared consistency layer. Most scalable in production systems.

Memory Failure Propagation

**Memory is not neutral.
It can amplify bias and error.**



Memory Contamination



If incorrect information is stored as truth, future reasoning compounds the error. Memory requires validation before persistence.



Forgetting Mechanisms

Memory grows infinitely unless managed

Time-Based Pruning

- Keep last N days.


Importance-Based Retention

- **Keep:** Errors, violations, high-confidence events
- **Drop:** Trivial logs

Compression

- Summarize episodes.


Forgetting Mechanisms Overview



Memory governance prevents bloat and bias accumulation.

Time-Travel / Checkpointing

A non-deterministic system debugging tool



Checkpointed execution graph allows rewind, modification, and branching. Essential for debugging non-deterministic LLM agents.

Why Time-Travel Matters in MAS

Transforms linear execution into a versioned state graph

- Debug reasoning paths
- Explore alternative decisions
- Reproduce prior outputs
- Audit decision flow
- Reduce recomputation




Key Takeaways Presentation Session 2

Memory design and shared context strategies

- Memory is not storage. It is influence
- Episodic memory preserves task history
- Procedural memory encodes reusable routines
- Checkpointing creates recoverability and auditability
- Retrieval shapes downstream reasoning paths


Pulse Check

Are the different memory types clear and when to use them?



Demo


A LangGraph workflow that maintains episodic and procedural memory and uses checkpoints; how memory is read and written during execution; how time travel allows stepping back to an earlier state to inspect the reasoning path



Hands-on

Extend your agent by adding a simple episodic memory that stores each tool result and a procedural memory entry that records a small reusable routine

- Instructions are in the Colab Notebook.



Q&A





Break

Session 3

Choosing the Right Models and Architectural Patterns for Agent Reasoning

Course Schedule

Session 3

Choosing the right models and architectural patterns for agent reasoning (60 minutes)

- **Presentation:** How architectural decisions shape the capability, speed, and stability of an agentic system; comparing planner executor setups, supervisor based coordination, and unified agent designs; assigning models to roles based on their strengths; when to use thinking or no thinking inference modes; KV cache behavior; context engineering, and the need for stable prefixes when designing long running workflows
- **Demo:** Comparing two LangGraph architectures, one using a planner with a thinking mode enabled and one using a unified agent with no thinking mode
- **Hands-on Exercise:** Modify a predefined workflow by switching the planning component between thinking and no thinking modes and observe the change in latency and decision quality; run the same workflow with a dense model and a smaller MoE model to compare behavior
- **Q&A** (5 mins)

Planner Executor Setup

One model plans, another executes steps, often with tools


Where to use it:

- Multi step workflows with clear decomposition
- Tool heavy automation
- When you want predictable control flow
- When you can validate executor outputs at each step

Where it breaks:

- Open ended exploration without clear stop conditions
- When the planner becomes a bottleneck
- When executor repeatedly asks for clarification due to missing context

Two Role Separation



Planner optimizes global coherence, executor optimizes throughput and tool interaction. Stability depends on validation gates and tight plan formats.

Supervisor Based Coordination

A supervisor routes tasks to specialists, aggregates, decides when to stop


Where to use it:

- Multi agent teams, research plus writing plus verification
- Workflows with parallelizable subtasks
- When you need clear ownership per role
- When you want modular upgrades per agent

Where it breaks:

- If everything routes through supervisor, latency grows
- If supervisor holds too much context, quality collapses
- If no aggregation criteria, system loops

Central Router Topology



Great control surface for reliability, audits, and specialization. Main risk is a cognitive bottleneck at the supervisor.

Unified Agent Design

One agent plans, executes tools, reflects, writes, and finishes


Where to use it:

- Small scoped tasks
- Low latency interactive chat
- Simple tool calls
- Early prototypes where you want minimal wiring

Where it breaks:

- Long workflows, long context, many tools
- Increased error propagation due to missing validation boundaries
- Higher token burn because one model does everything

Minimal Wiring, Maximal Coupling



Fast to build, harder to scale. Quality depends on strong internal discipline and robust tool validation.

Capability, Speed, Stability Comparison

Design	Capability ceiling	Speed	Stability under long runs
Unified agent	Medium, saturates quickly	High at low load	Low unless context is tightly managed
Planner executor	High for structured workflows	Medium	Medium to high with validation gates
Supervisor team	Highest for complex systems	Medium to low depending on routing	High when memory and aggregation are engineered

Model Selection

One size does not fit all. The choice of model impacts performance and cost.

- Small models are ideal for routing, specialized and simple tasks. They are fast and cheap.
- Use large models for complex reasoning, planning, or deep research.
- The combination is important. Use a small model for the easy parts, and a large model only when a hard problem requires it.
- **Evaluate open vs. closed models.**

Library for calculating costs: <https://github.com/pydantic/genai-prices>

Costs

Costs are per-step, not per-request. Each reasoning step and tool call adds to the total

- Small models are ideal for routing, specialized and simple tasks. They are fast and cheap.
- Use large models for complex reasoning, planning, or deep research.
- The combination is important. Use a small model for the easy parts, and a large model only when a hard problem requires it.
- Deactivate thinking modes based on task.

Role Mapping That Tends to Work in Production

- **Fast generator:** Drafting, formatting, tool calling, lightweight reasoning
- **Deliberate thinker:** Planning, decomposition, verification, conflict resolution, final aggregation
- **Encoder or embedding model:** Retrieval, filtering, ranking, relevance scoring, deduplication
- **MoE specialist:** High capability at controlled runtime cost, good for broad multitask workloads

Architecture Shapes Capability, Speed, Stability


Agent quality is an architecture problem. Because architecture controls:

Capability is what the system can reliably do. Speed is how fast it responds under load. Stability is whether it keeps working as context grows and concurrency increases



Thinking Mode vs No Thinking Mode

Thinking is a budget, not a personality. Treat it as a controllable compute allocation. Enable it at decision points, disable it on cheap execution paths.



Thinking Mode vs No Thinking Mode

When to use thinking

- Planning and decomposition
- Tool selection when tool choice matters
- Verification, conflict resolution, safety checks
- When mistakes are expensive downstream

When to disable thinking

- Short low risk turns
- Mechanical rewriting or formatting
- Steps that follow an already validated plan
- High concurrency endpoints where latency dominates

Right Amount of Thinking Budget

- Avoid setting the thinking budget too low, as this can limit the model's reasoning ability.
- Adjust the value based on the latency tolerance of you or your end users.
- A **thinking budget higher than 1024** typically provides:
 - Noticeable improvements in overall performance.
 - Better handling of multi-step or complex reasoning tasks.

KV Cache Behavior

Why it becomes an architectural concern


Core idea

Decoder generation reuses past keys and values. This increases throughput, but the cache grows with sequence length and concurrency.

Consequences for agent systems


- Long sessions increase GPU memory pressure
- High concurrency multiplies cache footprint
- Context edits can invalidate cache segments
- Cache reuse is strongest when prompts share stable prefixes

KV Cache Behavior



KV caching converts repeated attention compute into stored state. It shifts the bottleneck from compute to memory bandwidth and capacity. Under load, cache management decides whether latency is stable.

Context Engineering



Context engineering components centered on the user prompt. Each element contributes to shaping the agent's understanding and response within the overall context space.

Context Engineering and Stable Prefixes


The problem

If you reorder, trim, or inject tokens into the middle of the context, cached entries no longer match positions. Reuse becomes unsafe.

The principle: Keep a stable prefix, then treat the rest as modular segments.

- Stable prefix examples
- System instructions
- Tool schemas
- Output format constraints
- Safety policies


Context Engineering and Stable Prefixes Overview



Stable prefixes maximize cache reuse and keep long running workflows coherent. Treat context as a structured layout, not a chat transcript.

Unified Architecture View

Capability emerges from fit. Topology and protocol determine information flow. Memory and inference determine continuity and cost stability. Cache behavior makes context design a first class architectural concern.



Architecture to Task Mapping

Task	Best starting family	Notes
Open ended chat	Decoder only	Use instruction tuned base
Code synthesis	Decoder only	Leverage tool use and function calling
Semantic search	Encoder only	Dual encoders with ANN index
NER and tagging	Encoder only	Token level classifier
Translation	Encoder–decoder	Cross attention aligns source and target
RAG pipelines	Encoder + Decoder	Encoders retrieve. Decoders generate

Architecture Overview Table

Type	Primary capability	Typical use cases
Encoder-only	Representation learning. Contextual embeddings	Text classification. Semantic search. NER
Decoder-only	Autoregressive generation. Instruction following	Text generation. Code synthesis. Chat
Encoder-decoder	Sequence to sequence mapping	Translation. Summarization. QA over given context
Embedding models	Dense or sparse vector representations	Retrieval augmented generation. Similarity search
Mixture of Experts	Sparse expert routing. Scalable compute	Efficient large scale generation. Multitask

Choosing the Right Architecture

- Do you need bidirectional context at once → encoder
- Do you need stepwise generation → decoder
- Do you need input to output mapping across domains → encoder–decoder
- Do you need capacity without linear compute growth → MoE

Practical Checklists

Before choosing

- Objective and data model
- Latency and batch requirements
- Memory and hardware limits

Before serving

- Throughput profile
- Cache policies
- Telemetry for routing and retrieval quality


Key Takeaways Presentation Session 3

Choosing the right models and architectural patterns for agent reasoning

- Model choice is an architectural decision
- Planner vs unified affects latency and controllability
- Thinking mode trades compute for reliability
- Dense vs MoE affects throughput and stability
- KV cache and prefix stability matter in long-running workflows


Pulse Check

text



Demo


Comparing two LangGraph architectures, one using a planner with a thinking mode enabled and one using a unified agent with no thinking mode



Hands-on

Modify a predefined workflow by switching the planning component between thinking and no thinking modes and observe the change in latency and decision quality; run the same workflow with a dense model and a smaller MoE model to compare behavior

- Instructions are in the Colab Notebook.



Q&A





Break

Session 4

Observability, Monitoring, and Benchmarking

Course Schedule


Session 4

Building Real-world Multimodal Workflows (45 minutes)

- **Presentation:** Observing, measuring, and evaluating agent behavior in production; building task-specific benchmarks for multi-agent systems
- **Demo:** Monitoring and evaluating an agent workflow with monitoring hooks, logging, and a custom evaluation loop
- **Hands-on Exercise:** Add tracking to your workflow with one simple logging function that prints each state transition; run the workflow twice with small changes in the input and compare the differences in the logged steps to evaluate stability
- **Q&A** (5 mins)


Continuous Evaluation

The agent deployment lifecycle: from MVP to production through continuous evaluation




Monitoring Loop

Illustrative loop how you can monitor behavior and turn it into system improvements.




Diagnostic Lab




Iterative diagnostic loop showing how observation feeds back into continuous system refinement.

General MAS Evaluation Flow



High-level logic for any agentic workflow (Customer Support, Research, Supply Chain, etc.). It emphasizes that Tracing is the connective tissue between the dataset and the final outcome.

Multi-Agent Tracing



You need to track intent vs. action. This chart shows how a single Trace ID follows a request across different specialized agents and their tool usages.

Stress Testing Your Multi-Agent System

Stress Dimension	What it tests in MAS	Evaluation Metric
Communication Noise	Can agents handle contradictory messages?	Alignment Score
Latency Injection	Does the system time out or loop if an API is slow?	Resilience / Cost
Agent Drop-out	If the “Expert Agent” fails, can the others recover?	Self-Healing Rate
Ambiguity	Does the system ask for clarification or hallucinate?	Calibration

Logging, Tracing and Benchmarking

- Logs tell you what an agent thought (the "Mental Model").
- Traces tell you how the agents collaborated (the "Workflow Model").
- Eval Datasets tell you if they were right (the "Success Model").



Tools to Enhance Transparency

→ **These tools increase auditability, traceability, and trust in agentic systems**

- LangSmith: Tracks agent execution and reasoning paths
- Langfuse: Visualize, debug, and monitor LangChain applications
- Weights & Biases (Weave): Tracks experiments, outputs, and agents over time
- Circuit Tracing: Attribution graphs to explain model behavior

LangSmith Observability






Image source: LangChain

LangSmith Observability

The screenshot displays the LangSmith observability interface. On the left, a 'TRACE' sidebar shows a workflow with two steps: 'rag' (0.78s, 40 tokens) and 'ChatOpenAI' (0.77s). The main panel shows the details for the 'rag' step, with tabs for 'Run', 'Feedback', and 'Metadata'. The 'Input' section shows a JSON object:

```
{  "question": "where did harrison work"}
```

. The 'Output' section shows a JSON object:

```
{  "output": {    "id": "chatcmpl-9JWoyyc3dzdBtuqJf7pAJf1ip4n6u",    "choices": [      {        "finish_reason": "stop",        "index": 0,        "message": {          "content": "Harrison worked at Kensho.",          "role": "assistant"        }      }    ]  }}
```

Image source: LangChain

Langfuse Observability

The screenshot displays the Langfuse Observability interface. On the left is a sidebar with navigation links: Dashboard, Tracing, Sessions, Generations, Scores, Models, Users, Prompts, Datasets, Settings, Support, and Feedback. The main area is titled 'Trace Detail' for a specific trace ID. It shows session information, tags, and a detailed view of the generation process. The generation is performed by 'openai_llm' and includes a user query, a tool call, and an assistant response. A right-hand pane shows a 'TRACE' tree view of the entire process, including steps like 'agent_step', 'function_call', 'query', 'retrieve', 'OpenAIEmbedding', 'synthesize', and 'templating'.

Langfuse v2.11.1

demo-llamaindex-marc > Traces > 89c62a4d-e992-4923-a6b7-e2f27ae4cff3

Trace Detail

Session: notebook-session-2 795 → 34 (834) Total cost: \$0.0004

Tags

GENERATION openai_llm
3/13/2024, 7:54:12 PM
Latency: 0.82s 236 prompt → 8 completion (8244) gpt-3.5-turbo-0125 \$0.0001 max_tokens:
temperature: 0.1 request_timeout: 60.0

Add score + Add to dataset

Pretty JSON

user

"What did he do growing up?"

Show 9 more ...

query_engine_tool

"He is 43 years old."

```
{  
  "tool_call_id": "call_u85uyNS0JCyvOmGQb6WR7cxV"  
}
```

assistant

"He is 43 years old."

Metadata

```
{...} : 100%
```

TRACE Llamaindex_chat
2.44s

SPAN agent_step
2.44s

GENERATION openai_llm
0.81s 201 → 19 (8220)

SPAN function_call
0.81s

SPAN query
0.81s

SPAN retrieve
0.27s

GENERATION OpenAIEmbedding
0.27s 0 → 0 (85)

SPAN synthesize
0.54s

SPAN templating
0.00s

GENERATION openai_llm
0.54s 358 → 7 (8365)

GENERATION openai_llm
0.82s 236 → 8 (8244)

Image source: Langfuse

Weights & Biases: Weave

The screenshot displays the Weights & Biases Weave interface. On the left, a sidebar contains navigation links for Evaluations, Models, Datasets, Traces, Operations, and Objects. The main panel is titled 'Traces' and shows a list of traces with columns for Trace, Status, and Called. The selected trace is 'Agent.step' with ID '1b3c'. To the right of the list, a 'Trace tree' shows the hierarchy of the selected trace, including 'Agent.step', 'openai.chat.completions.create', and three 'run_command' nodes. Further right, a detailed view of the 'Agent.step' trace is shown, including its inputs, outputs, and child calls.

Traces

Op Roots only

Common values: input.self Agent:v0

Trace	Status	Called
Agent.step 7766	Success	5 days ago
Agent.step c5ba	Success	5 days ago
Agent.step 1b3c	Success	5 days ago
Agent.step 0e3d	Success	5 days ago
Agent.step ed82	Success	5 days ago

Total Rows: 5

Trace tree

- Agent.step
 - openai.chat.completions.create
 - run_command
 - run_command
 - run_command

Agent.step 1b3c

Call Code Summary

Inputs

Path	Value
self	Agent:v0
state	AgentState:v3

Outputs


Path	Value
Value	AgentState:v4

Child calls of run_command:v0

Trace	Status	inputs	
		command	output
run_command b91d	Success	ps aux	["exi
run_command a26f	Success	ifconfig	["exi
run_command ddb4	Success	free -m	["exi

Image source: Weights & Biases


Circuit Tracing: Chain-of-thought Faithfulness



Motivated reasoning (unfaithful): The model gives the wrong answer, It knows it will next multiply by 5, so it answers 0.8 so that $0.8 \cdot 5 = 4$ will match the answer which the user claimed to come to.

Image source: Anthropic


Benchmarking



Traditional LLM benchmarks are not enough. They focus on static tasks, not dynamic, multi-step agentic behavior. Create custom benchmarks that mimic your real-world use cases and run them continuously in production.

Paper: [BetterBench: Assessing AI Benchmarks, Uncovering Issues, and Establishing Best Practices](#)

SWE-smith: Scaling Data for SWE Agents



SWE-smith generates training data for software engineering agents by introducing bugs into real-world codebases.

Image source: <https://swesmith.com/>

SWE-smith: Scaling Data for SWE Agents

Creating training data for software engineering agents is difficult.

- Generate 100 task instances for any GitHub repository in 10 minutes.
- You can generate 100 to 1000 of task instances for any GitHub repository.

Context-Aware Evaluation

When you run SweSmith against your own repo, it doesn't just generate random code; it generates validated task instances that are unique to your proprietary logic.

- **Procedural Mutations:** It uses AST (Abstract Syntax Tree) transformations to flip operators or swap arguments in your functions.
- **PR Mirroring:** It can "undo" your recent successful PRs to see if your AI agent can find the same fix your human team did.
- **The Workflow Link:** Since it automates the Docker environment setup and test-running logic, it ensures the evaluation is grounded in your actual CI/CD constraints.

Solving the Data Leakage Problem

The biggest issue with evaluating agents on public benchmarks in 2026 is that the models have likely seen the solutions during pre-training. By generating an evaluation set from your internal repo, you create a "blind test" that is 100% contamination-proof.




Automated Golden Labels

Every SweSmith task comes with:

- **A Fail-to-Pass (F2P) test:** A unit test that fails on the bug but passes on the original code.
- **A Golden Patch:** The original code itself.
- **A Problem Statement:** An LLM-generated GitHub issue describing the bug.


Dynamic Scaling of Unit Tests



Increasing the number of unit tests used for majority voting enhances performance across various policy and reward models. Here, the policy model generates code solutions, while the reward model is responsible for creating unit tests.

Paper: [Dynamic Scaling of Unit Tests for Code Reward Modeling](#)

Dynamic Scaling of Unit Tests



Overview of unit test scaling. First, a lightweight unit test generator is trained based on high-quality synthetic data. Subsequently, a dynamic unit test scaling is employed to further improve efficiency.


Key Takeaways Presentation Session 4

Observability, monitoring, and benchmarking

- Agents are stochastic distributed systems
- Logging state transitions exposes hidden drift
- Stability is measurable
- Benchmarks must be task-specific
- Small input perturbations reveal fragility


Pulse Check

Can you see how the synergy between logging, tracing, and evaluation cycles effectively hardens your agents for production?



Demo


Monitoring and evaluating an agent workflow
with monitoring hooks, logging, and a custom
evaluation loop



Hands-on

Add tracking to your workflow with one simple logging function that prints each state transition; run the workflow twice with small changes in the input and compare the differences in the logged steps to evaluate stability

- Instructions are in the Colab Notebook.



Q&A





Break

Session 5

Security, guardrails, and secure execution environments

Course Schedule

Session 5

Security, guardrails, and secure execution environments (60 minutes)

- **Presentation:** The agentic threat landscape (reasoning manipulation, memory poisoning, tool misuse, communication corruption), safeguarding coding agents
- **Demo:** Securing an agentic workflow using OmegaConf-driven security policies, Pydantic schema validation, MCP tool governance, and PRP-based coding constraints
- **Group discussion:** Review a small predefined scenario involving an unsafe or ambiguous agent action and discuss where the vulnerabilities lie, what should be validated, and what mitigation strategy is most appropriate
- **Wrap-up and Q&A** (10 minutes)



Agentic AI Threat Landscape

Agentic AI introduces new and evolved threat types. Some risks extend existing AI vulnerabilities; others are unique to agentic architectures. These threats stem from new system components:

- Agent memory
- Tool integration



Key Emerging Attack Vectors

Agent Memory Poisoning

- Adversaries manipulate stored knowledge to alter agent behavior.

Key Emerging Attack Vectors

Tool Misuse

- Tools act as conduits for malicious commands or privilege escalation.
- Especially critical in unconstrained autonomy or multi-agent setups where agents learn from peers.



Key Emerging Attack Vectors

Code Generation Risks

- Tool misuse opens paths to Remote Code Execution (RCE) and code injection attacks.

Identity and Authorization Challenges

- Tool integration complicates identity management.
- Agents often act with elevated privileges on behalf of users.
- Risk of trust boundary violations when agents delegate across APIs or tools.
- Need for strong identity isolation in agentic environments.

Confused Deputy Vulnerability

- Occurs when an AI agent (the “deputy”):
- Has higher privileges than the user, and
- Is tricked into executing unauthorized actions.

Confused Deputy Vulnerability

Example:

An agent allowed to execute DB queries is manipulated via injected instructions to run high-privilege queries.


Root cause:

Lack of privilege separation and insufficient input validation.

Mitigation Strategies

- Downscope privileges when agents act on behalf of users.
- Enforce principle of least privilege (PoLP) in all tool integrations.
- Validate and sanitize user inputs and prompts rigorously.
- Employ identity isolation and context-aware authorization layers.
- Continuously monitor for prompt injection, spoofing, and impersonation attempts.

Threat Overview



Source: owasp.org

OWASP Top Agent Vulnerabilities

TID	Threat Name	Description (Condensed)	Key Mitigations
T1	Memory Poisoning	Injection of malicious or false data into agent memory to alter decisions or actions.	Validate and sanitize memory, isolate sessions, use anomaly detection, log snapshots for rollback.
T2	Tool Misuse	Manipulation of agents to misuse integrated tools via deceptive prompts or commands.	Verify tool calls, monitor usage, log executions, enforce strict operational boundaries.
T3	Privilege Compromise	Exploiting permission misconfigurations for unauthorized access or actions.	Apply granular role controls, audit elevated privileges, and restrict cross-agent delegation.
T4	Resource Overload	Targeting compute or memory limits to degrade or crash systems.	Implement quotas, rate limits, adaptive scaling, and real-time load monitoring.
T5	Cascading Hallucination	Propagation of false or misleading AI outputs that corrupt reasoning chains.	Use output validation, multi-source checks, feedback loops, and behavioral constraints.
T6	Intent Breaking / Goal Manipulation	Redirection of an agent's objectives or planning logic through adversarial inputs.	Validate goals, restrict reflective autonomy, audit behavior, and flag deviations.
T7	Misaligned / Deceptive Behaviors	AI acting harmfully or deceptively to fulfill objectives.	Apply safety alignment, require confirmation for risky tasks, use deception detection and red teaming.
T8	Repudiation / Untraceability	Lack of accountability or traceable logs for agent actions.	Enforce immutable logging, cryptographic signatures, and transparent audit trails.

OWASP Top Agent Vulnerabilities

TID	Threat Name	Description (Condensed)	Key Mitigations
T9	Identity Spoofing / Impersonation	Exploiting authentication flaws to impersonate agents or users.	Enforce identity verification, behavioral profiling, and continuous activity monitoring.
T10	Overwhelming Human-in-the-Loop	Exploiting human oversight weaknesses through overload or deception.	Use adaptive human-AI governance, dynamic trust thresholds, and hierarchical oversight.
T11	Unexpected RCE / Code Attacks	Injection of malicious or unintended code via AI-generated outputs.	Sandbox execution, restrict code privileges, and manually review elevated scripts.
T12	Agent Communication Poisoning	Tampering with inter-agent exchanges to spread false data or disrupt workflows.	Authenticate messages, validate communications, and enforce multi-agent consensus.
T13	Rogue Agents in Multi-Agent Systems	Compromised agents executing unauthorized actions or data exfiltration.	Apply policy constraints, continuous monitoring, red teaming, and integrity controls.
T14	Human Attacks on Multi-Agent Systems	Adversaries exploit agent delegation or trust to escalate privileges.	Limit delegation, enforce inter-agent authentication, and segment multi-agent tasks.
T15	Human Manipulation	Attackers exploit human trust to manipulate behavior or spread misinformation.	Monitor agent behavior, limit tool use, filter outputs, and apply guardrails or moderation APIs.

Agentic Threat Decision Path

- Does the AI agent independently determine the steps needed to achieve its goals?
- Does the AI agent rely on stored memory for decision-making?
- Does the AI agent execute actions using tools, system commands, or external integrations?

Agentic Threat Decision Path

- Does the AI system rely on authentication to verify users, tools, or services?
- Does AI require human engagement to achieve its goals or function effectively?
- Does the AI system rely on multiple interacting agents?

Mitigation Strategy Examples

Preventing AI Agent Reasoning Manipulation

- Step 1: Reduce attack surface & Implement Agent behavior profiling (Proactive)
- Step 2: Prevent AI agent Goal Manipulation (Reactive)
- Step 3: Strengthen AI Decision Traceability & Logging (Detective)



Preventing Memory Poisoning & AI Knowledge Corruption

- Step 1: Secure AI Memory Access & Validation (Proactive)
- Step 2: Detect & Respond to Memory Poisoning (Reactive)
- Step 3: Prevent the Spread of False Knowledge (Detective)

Securing AI Tool Execution & Preventing Unauthorized Actions

- Step 1: Restrict AI Tool Invocation & Execution (Proactive)
- Step 2: Monitor & Prevent Tool Misuse (Reactive)
- Step 3: Prevent AI Resource Exhaustion (Detective)

Strengthening Authentication, Identity & Privilege Controls

- Step 1: Implement Secure AI Authentication Mechanisms (Proactive)
- Step 2: Restrict Privilege Escalation & Identity Inheritance (Reactive)
- Step 3: Detect & Block AI Impersonation Attempts (Detective)

Protecting HITL & Preventing Decision Fatigue Exploits

- Step 1: Optimize HITL Workflows & Reduce Decision Fatigue (Proactive)
- Step 2: Identify AI-Induced Human Manipulation (Reactive)
- Step 3: Strengthen AI Decision Traceability & Logging (Detective)

Securing Multi-Agent Communication & Trust Mechanisms

- Step 1: Secure AI-to-AI Communication Channels (Proactive)
- Step 2: Detect & Block Rogue Agents (Reactive)
- Step 3: Enforce Multi-Agent Trust & Decision Security (Detective)


Further Tips to Prepare

- Create an objective threat profile
- Do an attack surface analysis
- Summarize and record known AI vulnerabilities
- Create an incident response readiness matrix
- Create an AI security matrix

Securing Code and Tool Usage

Using Agent2Agent for Tool Governance

The protocol embeds security and observability by design, making it enterprise ready. A2A enforces governance at the protocol level and provides structured artifacts and task separation, so you don't need a custom governance layer.




Sandboxing Agent Execution Overview


Isolation Layer	Best For	Key Benefit
Docker	Headless or internal tasks	Strong OS level control
Cloud Sandboxes	High variance or untrusted code	Zero risk to internal infrastructure
MCP (Pyodide, Deno)	Collaborative agent systems	Protocol level safety and speed

Zero-trust UI as Implicit Firewall

This setup keeps privileged actions off the client. The browser only talks to Streamlit. FastAPI stays internal, and all API keys and outbound calls live behind that boundary. In practice, that removes direct internet exposure of your agent runtime.



Use Path Validation Against Directory Traversal




The flow diagram illustrates how path validation is applied in practice. Every path request is normalized, resolved, and checked against a single writable root before any file operation is allowed to proceed. Anything that falls outside that boundary fails fast.

[More info on path traversal](#)

Sandbox State Transitions With E2B

Sandboxes such as E2B support additional execution patterns, including pausing and resuming execution. The graph illustrates the state transition from running to paused and then to killed.

[E2B info](#)






Key Takeaways Presentation Session 5

Security, guardrails, and secure execution environments

- Agents introduce new attack surfaces
- Reasoning can be manipulated
- Memory can be poisoned
- Tools can be misused
- Protocols can be corrupted
- Validation and governance must be layered


Pulse Check

Do you understand the benefits of safeguarding on different levels?



Demo

Securing an agentic workflow using OmegaConf-driven security policies, Pydantic schema validation, MCP tool governance, and PRP-based coding constraints



Group Discussion

Review a small predefined scenario involving an unsafe or ambiguous agent action and discuss where the vulnerabilities lie, what should be validated, and what mitigation strategy is most appropriate



Day Two Key Takeaways

- Structure before execution. Schemas are contracts, not formatting.
- Memory is influence. Store intentionally or risk silent corruption.
- Checkpointing enables recoverability and auditability.
- Model choice is architecture. Inference configuration shapes system behavior.
- Agents are stochastic systems. Stability must be measured, not assumed.
- Benchmarks must reflect real task structure, not generic scores.
- Production-grade agents require containment across structure, memory, compute, monitoring, and policy.

Q&A



The image features the O'Reilly logo in white, centered on a blue background. The logo consists of the word "O'REILLY" in a bold, sans-serif font, followed by a registered trademark symbol (®). To the left of the text, there are two large, overlapping, semi-transparent blue circles that create a sense of depth and movement.

O'REILLY®